

train

March 8, 2025

```
[12]: import torch
import numpy as np
from torch.utils.data import DataLoader, random_split, TensorDataset
import torch.nn as nn
import torch.optim as optim
from transformers import ViTForImageClassification, ViTConfig
import time
```

```
[13]: # Use CPU
device = "cpu"
print(f"Using device: {device}")
```

Using device: cpu

```
[14]: # Load and modify the ViT configuration
print("Initializing ViT model...")
config = ViTConfig.from_pretrained("google/vit-base-patch16-224")
config.num_channels = 1 # Change from RGB (3 channels) to Grayscale (1 channel)
config.image_size = 224 # Ensure model expects 224x224 input
```

Initializing ViT model...

```
[15]: # Initialize model with modified config
model = ViTForImageClassification(config)

# Adjust patch embeddings to accept single-channel input
model.vit.embeddings.patch_embeddings.projection = nn.Conv2d(
    in_channels=1, out_channels=config.hidden_size, kernel_size=16, stride=16
)

# Move model to CPU
model.to(device)
```

```
[15]: ViTForImageClassification(
  (vit): ViTModel(
    (embeddings): ViTEmbeddings(
      (patch_embeddings): ViTPatchEmbeddings(
        (projection): Conv2d(1, 768, kernel_size=(16, 16), stride=(16, 16))
      )
    )
  )
```

```

        (dropout): Dropout(p=0.0, inplace=False)
    )
    (encoder): ViTEncoder(
      (layer): ModuleList(
        (0-11): 12 x ViTLayer(
          (attention): ViTSdpaAttention(
            (attention): ViTSdpaSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.0, inplace=False)
            )
          (output): ViTSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.0, inplace=False)
          )
        )
      )
    (intermediate): ViTIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
      (intermediate_act_fn): GELUActivation()
    )
    (output): ViTOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (dropout): Dropout(p=0.0, inplace=False)
    )
    (layernorm_before): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
    (layernorm_after): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
  )
)
(layernorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)
(classifier): Linear(in_features=768, out_features=1000, bias=True)
)

```

```

[16]: # Function to pad images from (64, 72) to (224, 224)
def pad_images(images):
    padded = torch.zeros((images.shape[0], 1, 224, 224), dtype=torch.float32)
    padded[:, :, :64, :72] = images # Place original images in top-left corner
    return padded

# Load datasets
print("Loading datasets...")
data1 = np.load("Run355456_Dataset_jqkne.npy") # Shape: (N1, 64, 72)
data2 = np.load("Run357479_Dataset_iodic.npy") # Shape: (N2, 64, 72)

```

```

# Normalize data (important for convergence)
print("Normalizing datasets...")
data1 = data1 / data1.max()
data2 = data2 / data2.max()

# Assign labels
labels1 = np.zeros(len(data1), dtype=np.int64) # Class 0
labels2 = np.ones(len(data2), dtype=np.int64)  # Class 1

# Merge datasets
print("Merging datasets...")
data = np.concatenate((data1, data2), axis=0) # (Total_N, 64, 72)
labels = np.concatenate((labels1, labels2), axis=0) # (Total_N,)

# Convert to PyTorch tensors
print("Converting datasets to PyTorch tensors...")
data = torch.tensor(data, dtype=torch.float32).unsqueeze(1) # (N, 1, 64, 72)
data = pad_images(data) # Convert (1, 64, 72) → (1, 224, 224)
labels = torch.tensor(labels, dtype=torch.long)

```

Loading datasets...

Normalizing datasets...

Merging datasets...

Converting datasets to PyTorch tensors...

```

[17]: # Train/Val/Test split (70-20-10)
train_size = int(0.7 * len(data))
val_size = int(0.2 * len(data))
test_size = len(data) - train_size - val_size

print(f"Dataset split: {train_size} train, {val_size} val, {test_size} test.")
train_data, val_data, test_data = random_split(TensorDataset(data, labels),
    ↪ [train_size, val_size, test_size])

# Initialize DataLoaders with smaller batch size for CPU
batch_size = 16 # Reduced batch size for CPU
print(f"Initializing DataLoaders with batch size {batch_size}...")
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True,
    ↪ num_workers=0)
val_loader = DataLoader(val_data, batch_size=batch_size, num_workers=0)
test_loader = DataLoader(test_data, batch_size=batch_size, num_workers=0)

```

Dataset split: 14000 train, 4000 val, 2000 test.

Initializing DataLoaders with batch size 16...

```

[18]: # Loss and Optimizer
print("Setting up optimizer and loss function...")

```

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=2e-4, weight_decay=1e-4)

# Learning rate scheduler
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer, mode='max', factor=0.5, patience=2, verbose=True
)

```

Setting up optimizer and loss function...

```

[19]: # Training parameters
epochs = 5
best_val_acc = 0.0
best_train_acc = 0.0

```

```

[20]: # Helper function for validation
def validate(model, val_loader):
    model.eval()
    val_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in val_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images).logits
            loss = criterion(outputs, labels)

            val_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    return val_loss / len(val_loader), correct / total

```

```

[24]: print("\n--- Training Starts ---")
for epoch in range(epochs):
    start_time = time.time()
    print(f"\n Epoch {epoch+1}/{epochs} starts...")

    # Training phase
    model.train()
    total_loss = 0.0
    correct = 0
    total = 0

    for batch_idx, (images, labels) in enumerate(train_loader):
        images, labels = images.to(device), labels.to(device)

```

```

    # Forward pass
    outputs = model(images).logits
    loss = criterion(outputs, labels)

    # Backward pass and optimize
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    # Statistics
    total_loss += loss.item()
    _, predicted = torch.max(outputs.data, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

    # Print progress
    if (batch_idx + 1) % 10 == 0 or batch_idx + 1 == len(train_loader):
        print(f" Batch {batch_idx + 1}/{len(train_loader)} | "
              f"Loss: {total_loss/(batch_idx+1):.4f} | "
              f"Acc: {correct/total:.4f}")

train_loss = total_loss / len(train_loader)
train_acc = correct / total

# Validation phase
val_loss, val_acc = validate(model, val_loader)

# Update scheduler based on validation accuracy
scheduler.step(val_acc)

# Save best model
if val_acc > best_val_acc or train_acc > best_train_acc:
    best_val_acc = val_acc
    best_train_acc = train_acc
    torch.save(model.state_dict(), "best_vit_model_jupyternotebook.pth")
    print(f" New best model saved with validation accuracy: {val_acc:.4f}")

# Print epoch summary
epoch_time = time.time() - start_time
print(f" Epoch {epoch+1} completed in {epoch_time:.1f}s | "
      f"Train Loss: {train_loss:.4f} | Train Acc: {train_acc:.4f} | "
      f"Val Loss: {val_loss:.4f} | Val Acc: {val_acc:.4f}")

```

--- Training Starts ---

Epoch 1/5 starts...

Batch 10/875	Loss: 0.7614	Acc: 0.4375
Batch 20/875	Loss: 0.7382	Acc: 0.4813
Batch 30/875	Loss: 0.7245	Acc: 0.5021
Batch 40/875	Loss: 0.6743	Acc: 0.5750
Batch 50/875	Loss: 0.5607	Acc: 0.6600
Batch 60/875	Loss: 0.4675	Acc: 0.7167
Batch 70/875	Loss: 0.4065	Acc: 0.7562
Batch 80/875	Loss: 0.3717	Acc: 0.7805
Batch 90/875	Loss: 0.3440	Acc: 0.8028
Batch 100/875	Loss: 0.3125	Acc: 0.8213
Batch 110/875	Loss: 0.2857	Acc: 0.8375
Batch 120/875	Loss: 0.2623	Acc: 0.8510
Batch 130/875	Loss: 0.2428	Acc: 0.8620
Batch 140/875	Loss: 0.2289	Acc: 0.8710
Batch 150/875	Loss: 0.2152	Acc: 0.8792
Batch 160/875	Loss: 0.2023	Acc: 0.8867
Batch 170/875	Loss: 0.1904	Acc: 0.8934
Batch 180/875	Loss: 0.1799	Acc: 0.8993
Batch 190/875	Loss: 0.1704	Acc: 0.9046
Batch 200/875	Loss: 0.1619	Acc: 0.9094
Batch 210/875	Loss: 0.1567	Acc: 0.9134
Batch 220/875	Loss: 0.1498	Acc: 0.9173
Batch 230/875	Loss: 0.1434	Acc: 0.9209
Batch 240/875	Loss: 0.1375	Acc: 0.9242
Batch 250/875	Loss: 0.1320	Acc: 0.9273
Batch 260/875	Loss: 0.1269	Acc: 0.9300
Batch 270/875	Loss: 0.1223	Acc: 0.9326
Batch 280/875	Loss: 0.1179	Acc: 0.9350
Batch 290/875	Loss: 0.1139	Acc: 0.9373
Batch 300/875	Loss: 0.1101	Acc: 0.9394
Batch 310/875	Loss: 0.1066	Acc: 0.9413
Batch 320/875	Loss: 0.1032	Acc: 0.9432
Batch 330/875	Loss: 0.1001	Acc: 0.9449
Batch 340/875	Loss: 0.0972	Acc: 0.9465
Batch 350/875	Loss: 0.0944	Acc: 0.9480
Batch 360/875	Loss: 0.0918	Acc: 0.9495
Batch 370/875	Loss: 0.0893	Acc: 0.9508
Batch 380/875	Loss: 0.0870	Acc: 0.9521
Batch 390/875	Loss: 0.0854	Acc: 0.9532
Batch 400/875	Loss: 0.0833	Acc: 0.9544
Batch 410/875	Loss: 0.0813	Acc: 0.9555
Batch 420/875	Loss: 0.0794	Acc: 0.9565
Batch 430/875	Loss: 0.0775	Acc: 0.9576
Batch 440/875	Loss: 0.0758	Acc: 0.9585
Batch 450/875	Loss: 0.0741	Acc: 0.9594
Batch 460/875	Loss: 0.0725	Acc: 0.9603
Batch 470/875	Loss: 0.0710	Acc: 0.9612

Batch 480/875 | Loss: 0.0695 | Acc: 0.9620
Batch 490/875 | Loss: 0.0681 | Acc: 0.9628
Batch 500/875 | Loss: 0.0667 | Acc: 0.9635
Batch 510/875 | Loss: 0.0654 | Acc: 0.9642
Batch 520/875 | Loss: 0.0642 | Acc: 0.9649
Batch 530/875 | Loss: 0.0630 | Acc: 0.9656
Batch 540/875 | Loss: 0.0618 | Acc: 0.9662
Batch 550/875 | Loss: 0.0607 | Acc: 0.9668
Batch 560/875 | Loss: 0.0596 | Acc: 0.9674
Batch 570/875 | Loss: 0.0586 | Acc: 0.9680
Batch 580/875 | Loss: 0.0575 | Acc: 0.9685
Batch 590/875 | Loss: 0.0566 | Acc: 0.9691
Batch 600/875 | Loss: 0.0556 | Acc: 0.9696
Batch 610/875 | Loss: 0.0547 | Acc: 0.9701
Batch 620/875 | Loss: 0.0538 | Acc: 0.9706
Batch 630/875 | Loss: 0.0530 | Acc: 0.9710
Batch 640/875 | Loss: 0.0522 | Acc: 0.9715
Batch 650/875 | Loss: 0.0514 | Acc: 0.9719
Batch 660/875 | Loss: 0.0506 | Acc: 0.9723
Batch 670/875 | Loss: 0.0498 | Acc: 0.9728
Batch 680/875 | Loss: 0.0491 | Acc: 0.9732
Batch 690/875 | Loss: 0.0484 | Acc: 0.9736
Batch 700/875 | Loss: 0.0477 | Acc: 0.9739
Batch 710/875 | Loss: 0.0470 | Acc: 0.9743
Batch 720/875 | Loss: 0.0464 | Acc: 0.9747
Batch 730/875 | Loss: 0.0458 | Acc: 0.9750
Batch 740/875 | Loss: 0.0451 | Acc: 0.9753
Batch 750/875 | Loss: 0.0445 | Acc: 0.9757
Batch 760/875 | Loss: 0.0440 | Acc: 0.9760
Batch 770/875 | Loss: 0.0434 | Acc: 0.9763
Batch 780/875 | Loss: 0.0428 | Acc: 0.9766
Batch 790/875 | Loss: 0.0423 | Acc: 0.9769
Batch 800/875 | Loss: 0.0418 | Acc: 0.9772
Batch 810/875 | Loss: 0.0412 | Acc: 0.9775
Batch 820/875 | Loss: 0.0407 | Acc: 0.9777
Batch 830/875 | Loss: 0.0403 | Acc: 0.9780
Batch 840/875 | Loss: 0.0398 | Acc: 0.9783
Batch 850/875 | Loss: 0.0393 | Acc: 0.9785
Batch 860/875 | Loss: 0.0389 | Acc: 0.9788
Batch 870/875 | Loss: 0.0384 | Acc: 0.9790
Batch 875/875 | Loss: 0.0382 | Acc: 0.9791

New best model saved with validation accuracy: 1.0000

Epoch 1 completed in 2399.9s | Train Loss: 0.0382 | Train Acc: 0.9791 | Val
Loss: 0.0001 | Val Acc: 1.0000

Epoch 2/5 starts...

Batch 10/875 | Loss: 0.0001 | Acc: 1.0000
Batch 20/875 | Loss: 0.0001 | Acc: 1.0000

Batch 510/875 | Loss: 0.0001 | Acc: 1.0000
Batch 520/875 | Loss: 0.0001 | Acc: 1.0000
Batch 530/875 | Loss: 0.0001 | Acc: 1.0000
Batch 540/875 | Loss: 0.0001 | Acc: 1.0000
Batch 550/875 | Loss: 0.0001 | Acc: 1.0000
Batch 560/875 | Loss: 0.0001 | Acc: 1.0000
Batch 570/875 | Loss: 0.0001 | Acc: 1.0000
Batch 580/875 | Loss: 0.0001 | Acc: 1.0000
Batch 590/875 | Loss: 0.0001 | Acc: 1.0000
Batch 600/875 | Loss: 0.0001 | Acc: 1.0000
Batch 610/875 | Loss: 0.0001 | Acc: 1.0000
Batch 620/875 | Loss: 0.0001 | Acc: 1.0000
Batch 630/875 | Loss: 0.0001 | Acc: 1.0000
Batch 640/875 | Loss: 0.0001 | Acc: 1.0000
Batch 650/875 | Loss: 0.0001 | Acc: 1.0000
Batch 660/875 | Loss: 0.0001 | Acc: 1.0000
Batch 670/875 | Loss: 0.0001 | Acc: 1.0000
Batch 680/875 | Loss: 0.0001 | Acc: 1.0000
Batch 690/875 | Loss: 0.0001 | Acc: 1.0000
Batch 700/875 | Loss: 0.0001 | Acc: 1.0000
Batch 710/875 | Loss: 0.0001 | Acc: 1.0000
Batch 720/875 | Loss: 0.0001 | Acc: 1.0000
Batch 730/875 | Loss: 0.0001 | Acc: 1.0000
Batch 740/875 | Loss: 0.0001 | Acc: 1.0000
Batch 750/875 | Loss: 0.0001 | Acc: 1.0000
Batch 760/875 | Loss: 0.0001 | Acc: 1.0000
Batch 770/875 | Loss: 0.0001 | Acc: 1.0000
Batch 780/875 | Loss: 0.0001 | Acc: 1.0000
Batch 790/875 | Loss: 0.0001 | Acc: 1.0000
Batch 800/875 | Loss: 0.0001 | Acc: 1.0000
Batch 810/875 | Loss: 0.0001 | Acc: 1.0000
Batch 820/875 | Loss: 0.0001 | Acc: 1.0000
Batch 830/875 | Loss: 0.0001 | Acc: 1.0000
Batch 840/875 | Loss: 0.0001 | Acc: 1.0000
Batch 850/875 | Loss: 0.0001 | Acc: 1.0000
Batch 860/875 | Loss: 0.0001 | Acc: 1.0000
Batch 870/875 | Loss: 0.0001 | Acc: 1.0000
Batch 875/875 | Loss: 0.0001 | Acc: 1.0000

New best model saved with validation accuracy: 1.0000

Epoch 2 completed in 2600.9s | Train Loss: 0.0001 | Train Acc: 1.0000 | Val
Loss: 0.0000 | Val Acc: 1.0000

Epoch 3/5 starts...

Batch 10/875 | Loss: 0.0000 | Acc: 1.0000
Batch 20/875 | Loss: 0.0000 | Acc: 1.0000
Batch 30/875 | Loss: 0.0000 | Acc: 1.0000
Batch 40/875 | Loss: 0.0000 | Acc: 1.0000
Batch 50/875 | Loss: 0.0000 | Acc: 1.0000


```
Batch 540/875 | Loss: 0.0000 | Acc: 1.0000
Batch 550/875 | Loss: 0.0000 | Acc: 1.0000
Batch 560/875 | Loss: 0.0000 | Acc: 1.0000
Batch 570/875 | Loss: 0.0000 | Acc: 1.0000
Batch 580/875 | Loss: 0.0000 | Acc: 1.0000
Batch 590/875 | Loss: 0.0000 | Acc: 1.0000
Batch 600/875 | Loss: 0.0000 | Acc: 1.0000
Batch 610/875 | Loss: 0.0000 | Acc: 1.0000
Batch 620/875 | Loss: 0.0000 | Acc: 1.0000
Batch 630/875 | Loss: 0.0000 | Acc: 1.0000
Batch 640/875 | Loss: 0.0000 | Acc: 1.0000
Batch 650/875 | Loss: 0.0000 | Acc: 1.0000
Batch 660/875 | Loss: 0.0000 | Acc: 1.0000
Batch 670/875 | Loss: 0.0000 | Acc: 1.0000
Batch 680/875 | Loss: 0.0000 | Acc: 1.0000
Batch 690/875 | Loss: 0.0000 | Acc: 1.0000
Batch 700/875 | Loss: 0.0000 | Acc: 1.0000
Batch 710/875 | Loss: 0.0000 | Acc: 1.0000
Batch 720/875 | Loss: 0.0000 | Acc: 1.0000
Batch 730/875 | Loss: 0.0000 | Acc: 1.0000
Batch 740/875 | Loss: 0.0000 | Acc: 1.0000
Batch 750/875 | Loss: 0.0000 | Acc: 1.0000
Batch 760/875 | Loss: 0.0000 | Acc: 1.0000
Batch 770/875 | Loss: 0.0000 | Acc: 1.0000
Batch 780/875 | Loss: 0.0000 | Acc: 1.0000
Batch 790/875 | Loss: 0.0000 | Acc: 1.0000
Batch 800/875 | Loss: 0.0000 | Acc: 1.0000
Batch 810/875 | Loss: 0.0000 | Acc: 1.0000
Batch 820/875 | Loss: 0.0000 | Acc: 1.0000
Batch 830/875 | Loss: 0.0000 | Acc: 1.0000
Batch 840/875 | Loss: 0.0000 | Acc: 1.0000
Batch 850/875 | Loss: 0.0000 | Acc: 1.0000
Batch 860/875 | Loss: 0.0000 | Acc: 1.0000
Batch 870/875 | Loss: 0.0000 | Acc: 1.0000
Batch 875/875 | Loss: 0.0000 | Acc: 1.0000
Epoch 3 completed in 2807.7s | Train Loss: 0.0000 | Train Acc: 1.0000 | Val
Loss: 0.0000 | Val Acc: 1.0000
```

Epoch 4/5 starts...

```
Batch 10/875 | Loss: 0.0000 | Acc: 1.0000
Batch 20/875 | Loss: 0.0000 | Acc: 1.0000
Batch 30/875 | Loss: 0.0000 | Acc: 1.0000
Batch 40/875 | Loss: 0.0000 | Acc: 1.0000
Batch 50/875 | Loss: 0.0000 | Acc: 1.0000
Batch 60/875 | Loss: 0.0000 | Acc: 1.0000
Batch 70/875 | Loss: 0.0000 | Acc: 1.0000
Batch 80/875 | Loss: 0.0000 | Acc: 1.0000
Batch 90/875 | Loss: 0.0000 | Acc: 1.0000
```



```
Batch 580/875 | Loss: 0.0000 | Acc: 1.0000
Batch 590/875 | Loss: 0.0000 | Acc: 1.0000
Batch 600/875 | Loss: 0.0000 | Acc: 1.0000
Batch 610/875 | Loss: 0.0000 | Acc: 1.0000
Batch 620/875 | Loss: 0.0000 | Acc: 1.0000
Batch 630/875 | Loss: 0.0000 | Acc: 1.0000
Batch 640/875 | Loss: 0.0000 | Acc: 1.0000
Batch 650/875 | Loss: 0.0000 | Acc: 1.0000
Batch 660/875 | Loss: 0.0000 | Acc: 1.0000
Batch 670/875 | Loss: 0.0000 | Acc: 1.0000
Batch 680/875 | Loss: 0.0000 | Acc: 1.0000
Batch 690/875 | Loss: 0.0000 | Acc: 1.0000
Batch 700/875 | Loss: 0.0000 | Acc: 1.0000
Batch 710/875 | Loss: 0.0000 | Acc: 1.0000
Batch 720/875 | Loss: 0.0000 | Acc: 1.0000
Batch 730/875 | Loss: 0.0000 | Acc: 1.0000
Batch 740/875 | Loss: 0.0000 | Acc: 1.0000
Batch 750/875 | Loss: 0.0000 | Acc: 1.0000
Batch 760/875 | Loss: 0.0000 | Acc: 1.0000
Batch 770/875 | Loss: 0.0000 | Acc: 1.0000
Batch 780/875 | Loss: 0.0000 | Acc: 1.0000
Batch 790/875 | Loss: 0.0000 | Acc: 1.0000
Batch 800/875 | Loss: 0.0000 | Acc: 1.0000
Batch 810/875 | Loss: 0.0000 | Acc: 1.0000
Batch 820/875 | Loss: 0.0000 | Acc: 1.0000
Batch 830/875 | Loss: 0.0000 | Acc: 1.0000
Batch 840/875 | Loss: 0.0000 | Acc: 1.0000
Batch 850/875 | Loss: 0.0000 | Acc: 1.0000
Batch 860/875 | Loss: 0.0000 | Acc: 1.0000
Batch 870/875 | Loss: 0.0000 | Acc: 1.0000
Batch 875/875 | Loss: 0.0000 | Acc: 1.0000
Epoch 4 completed in 2210.7s | Train Loss: 0.0000 | Train Acc: 1.0000 | Val
Loss: 0.0000 | Val Acc: 1.0000
```

Epoch 5/5 starts...

```
Batch 10/875 | Loss: 0.0000 | Acc: 1.0000
Batch 20/875 | Loss: 0.0000 | Acc: 1.0000
Batch 30/875 | Loss: 0.0000 | Acc: 1.0000
Batch 40/875 | Loss: 0.0000 | Acc: 1.0000
Batch 50/875 | Loss: 0.0000 | Acc: 1.0000
Batch 60/875 | Loss: 0.0000 | Acc: 1.0000
Batch 70/875 | Loss: 0.0000 | Acc: 1.0000
Batch 80/875 | Loss: 0.0000 | Acc: 1.0000
Batch 90/875 | Loss: 0.0000 | Acc: 1.0000
Batch 100/875 | Loss: 0.0000 | Acc: 1.0000
Batch 110/875 | Loss: 0.0000 | Acc: 1.0000
Batch 120/875 | Loss: 0.0000 | Acc: 1.0000
Batch 130/875 | Loss: 0.0000 | Acc: 1.0000
```



```

Batch 620/875 | Loss: 0.0000 | Acc: 1.0000
Batch 630/875 | Loss: 0.0000 | Acc: 1.0000
Batch 640/875 | Loss: 0.0000 | Acc: 1.0000
Batch 650/875 | Loss: 0.0000 | Acc: 1.0000
Batch 660/875 | Loss: 0.0000 | Acc: 1.0000
Batch 670/875 | Loss: 0.0000 | Acc: 1.0000
Batch 680/875 | Loss: 0.0000 | Acc: 1.0000
Batch 690/875 | Loss: 0.0000 | Acc: 1.0000
Batch 700/875 | Loss: 0.0000 | Acc: 1.0000
Batch 710/875 | Loss: 0.0000 | Acc: 1.0000
Batch 720/875 | Loss: 0.0000 | Acc: 1.0000
Batch 730/875 | Loss: 0.0000 | Acc: 1.0000
Batch 740/875 | Loss: 0.0000 | Acc: 1.0000
Batch 750/875 | Loss: 0.0000 | Acc: 1.0000
Batch 760/875 | Loss: 0.0000 | Acc: 1.0000
Batch 770/875 | Loss: 0.0000 | Acc: 1.0000
Batch 780/875 | Loss: 0.0000 | Acc: 1.0000
Batch 790/875 | Loss: 0.0000 | Acc: 1.0000
Batch 800/875 | Loss: 0.0000 | Acc: 1.0000
Batch 810/875 | Loss: 0.0000 | Acc: 1.0000
Batch 820/875 | Loss: 0.0000 | Acc: 1.0000
Batch 830/875 | Loss: 0.0000 | Acc: 1.0000
Batch 840/875 | Loss: 0.0000 | Acc: 1.0000
Batch 850/875 | Loss: 0.0000 | Acc: 1.0000
Batch 860/875 | Loss: 0.0000 | Acc: 1.0000
Batch 870/875 | Loss: 0.0000 | Acc: 1.0000
Batch 875/875 | Loss: 0.0000 | Acc: 1.0000
Epoch 5 completed in 2348.5s | Train Loss: 0.0000 | Train Acc: 1.0000 | Val
Loss: 0.0000 | Val Acc: 1.0000

```

```

[28]: # Load best model for final evaluation
print("\n--- Final Evaluation ---")
model.load_state_dict(torch.load("best_vit_model_jupyternotebook.pth"))
model.eval()

```

--- Final Evaluation ---

```

[28]: ViTForImageClassification(
  (vit): ViTModel(
    (embeddings): ViTEmbeddings(
      (patch_embeddings): ViTPatchEmbeddings(
        (projection): Conv2d(1, 768, kernel_size=(16, 16), stride=(16, 16))
      )
      (dropout): Dropout(p=0.0, inplace=False)
    )
    (encoder): ViTEncoder(
      (layer): ModuleList(

```

```

(0-11): 12 x ViTLayer(
  (attention): ViTSdpaAttention(
    (attention): ViTSdpaSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.0, inplace=False)
    )
    (output): ViTSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.0, inplace=False)
    )
  )
  (intermediate): ViTIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
    (intermediate_act_fn): GELUActivation()
  )
  (output): ViTOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (dropout): Dropout(p=0.0, inplace=False)
  )
  (layernorm_before): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
  (layernorm_after): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
)
)
)
  (layernorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)
  (classifier): Linear(in_features=768, out_features=1000, bias=True)
)

```

```

[29]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Initialize test loss and tracking variables
test_loss = 0.0
correct = 0
total = 0
all_labels = []
all_probs = []

with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)

```



```

# Forward pass
outputs = model(images).logits
loss = criterion(outputs, labels)

# Accumulate test loss
test_loss += loss.item()

# Get predictions
probs = torch.softmax(outputs, dim=1)[: , 1] # Probability of class 1
_, predicted = torch.max(outputs.data, 1)

# Store values for evaluation
all_labels.extend(labels.cpu().numpy())
all_probs.extend(probs.cpu().numpy())

# Compute accuracy
total += labels.size(0)
correct += (predicted == labels).sum().item()

# Compute final loss and accuracy
test_loss = test_loss / len(test_loader)
test_acc = correct / total

# Compute ROC curve and AUC
fpr, tpr, _ = roc_curve(all_labels, all_probs)
roc_auc = auc(fpr, tpr)

# Print final test results
print(f"\n Final Test Results | Loss: {test_loss:.4f} | Accuracy: {test_acc:.4f} | AUC: {roc_auc:.4f}")

# Plot ROC Curve
plt.figure()
plt.plot(fpr, tpr, color="blue", lw=2, label=f"ROC curve (AUC = {roc_auc:.4f})")
plt.plot([0, 1], [0, 1], color="gray", linestyle="--") # Diagonal reference line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristic (ROC) Curve")
plt.legend(loc="lower right")
plt.show()

```

Final Test Results | Loss: 0.0000 | Accuracy: 1.0000 | AUC: 1.0000

